

# Lecture 10: External Merge Sort

# Today's Lecture

1. External merge sort
2. External merge sort on larger files
3. Optimizations for sorting

# 1. External Merge Sort

## Recap: External Merge Algorithm

- Suppose we want to merge two **sorted** files both much larger than main memory (i.e. the buffer)
- We can use the **external merge algorithm** to merge files of ***arbitrary length*** in  **$2*(N+M)$**  IO operations with only **3 buffer pages!**

Our first example of an “IO aware”  
algorithm / cost model

# Why are Sort Algorithms Important?

- Data requested from DB in sorted order is **extremely common**
  - e.g., find students in increasing GPA order
- **Why not just use quicksort in main memory??**
  - What about if we need to sort 1TB of data with 1GB of RAM...

A classic problem in computer science!

## More reasons to sort...


- Sorting useful for eliminating *duplicate copies* in a collection of records (Why?)
- Sorting is first step in *bulk loading* B+ tree index.
- *Sort-merge* join algorithm involves sorting

Coming up...

Next lecture

# Do people care?

<http://sortbenchmark.org>



Sort benchmark bears his name

# So how do we sort big files?


1. Split into chunks small enough to **sort in memory (“runs”)**
2. **Merge** pairs (or groups) of runs ***using the external merge algorithm***
3. **Keep merging** the resulting runs (***each time = a “pass”***) until left with one sorted file!

# External Merge Sort Algorithm

Example:

- 3 Buffer pages
- 6-page file

Orange file  
= unsorted




1. Split into chunks small enough to **sort in memory**

# External Merge Sort Algorithm

Example:

- 3 Buffer pages
- 6-page file

Orange file  
= unsorted




1. Split into chunks small enough to **sort in memory**

# External Merge Sort Algorithm

Example:

- 3 Buffer pages
- 6-page file

Orange file  
= unsorted




1. Split into chunks small enough to **sort in memory**

# External Merge Sort Algorithm

Example:

- 3 Buffer pages
- 6-page file

Orange file  
= unsorted




1. Split into chunks small enough to **sort in memory**

# External Merge Sort Algorithm

Example:

- 3 Buffer pages
- 6-page file

Each sorted  
file is a  
called a *run*




1. Split into chunks small enough to **sort in memory**

# External Merge Sort Algorithm

Example:

- 3 Buffer pages
- 6-page file



2. Now just run the **external merge** algorithm & we're done!


# Calculating IO Cost

For 3 buffer pages, 6 page file:

1. Split into **two 3-page files** and **sort in memory**
  1. =  $1 R + 1 W$  for each file =  $2 * (3 + 3) = 12$  IO operations
2. Merge each pair of sorted chunks ***using the external merge algorithm***
  1. =  $2 * (3 + 3) = 12$  IO operations
3. Total cost = **24** IO


## 2. External Merge Sort on larger files

# Running External Merge Sort on Larger Files



Assume we still  
only have 3 buffer  
pages (*Buffer not  
pictured*)


# Running External Merge Sort on Larger Files



1. Split into files small enough to sort in buffer...

Assume we still only have 3 buffer pages (*Buffer not pictured*)

# Running External Merge Sort on Larger Files




1. Split into files small enough to sort in buffer... and sort

Call each of these sorted files a *run*

Assume we still only have 3 buffer pages (*Buffer not pictured*)


# Running External Merge Sort on Larger Files



Assume we still  
only have 3 buffer  
pages (*Buffer not  
pictured*)

2. Now merge  
pairs of (sorted)  
files... **the  
resulting files  
will be sorted!**

# Running External Merge Sort on Larger Files




Assume we still  
only have 3 buffer  
pages (*Buffer not  
pictured*)

3. And repeat...

Call each of these  
steps a *pass*

# Running External Merge Sort on Larger Files




4. And repeat!

# Simplified 3-page Buffer Version

Assume for simplicity that we split an  $N$ -page file into  $N$  single-page ***runs*** and sort these; then:

- First pass: Merge  **$N/2$  pairs of runs** each of length **1 page**
- Second pass: Merge  **$N/4$  pairs of runs** each of length **2 pages**
- In general, for  **$N$  pages**, we do  **$\lceil \log_2 N \rceil$  passes**
  - +1 for the initial split & sort
- Each pass involves reading in & writing out all the pages =  **$2N$  IO**



→  $2N * (\lceil \log_2 N \rceil + 1)$  total IO cost!

# 3. Optimizations for sorting

# Using $B+1$ buffer pages to reduce # of passes

Suppose we have  $B+1$  buffer pages now; we can:

## 1. Increase length of initial runs. Sort $B+1$ at a time!

At the beginning, we can split the  $N$  pages into runs of length  $B+1$  and sort these in memory

IO Cost:

$$2N(\lceil \log_2 N \rceil + 1)$$



$$2N\left(\left\lceil \log_2 \frac{N}{B+1} \right\rceil + 1\right)$$

Starting with runs  
of length 1

Starting with runs of  
length  $B+1$

# Using $B+1$ buffer pages to reduce # of passes

Suppose we have  $B+1$  buffer pages now; we can:

## 2. Perform a $B$ -way merge.

On each pass, we can merge groups of  $B$  runs at a time (vs. merging pairs of runs)!

IO Cost:

$$2N(\lceil \log_2 N \rceil + 1) \rightarrow 2N\left(\left\lceil \log_2 \frac{N}{B+1} \right\rceil + 1\right) \rightarrow 2N\left(\left\lceil \log_B \frac{N}{B+1} \right\rceil + 1\right)$$

Starting with runs  
of length 1

Starting with runs of  
length  $B+1$

Performing  $B$ -way  
merges


# Rewriting

## Rewriting for even longer initial runs

- With  $B+1$  buffer pages, we can now start with  **$B+1$ -length initial runs** (and use  **$B$ -way merges**) to get  $2N(\left\lceil \log_B \frac{N}{B+1} \right\rceil + 1)$  IO cost...
- Can we reduce this cost more by getting even longer initial runs?
- Use **repacking**- produce longer initial runs by “merging” in buffer as we sort at initial stage

# Repacking Example: 3 page buffer


- Start with unsorted single input file, and load 2 pages



# Repacking Example: 3 page buffer


- Take the minimum two values, and put in output page

Also keep track of max (last) value in current run...




# Repacking Example: 3 page buffer

- Next, *repack*




# Repacking Example: 3 page buffer

- Next, **repack**, then load another page and continue!




## Repacking Example: 3 page buffer

- Now, however, ***the smallest values are less than the largest (last) in the sorted run...***




## Repacking Example: 3 page buffer

- Now, however, ***the smallest values are less than the largest (last) in the sorted run...***




## Repacking Example: 3 page buffer

- Now, however, ***the smallest values are less than the largest (last) in the sorted run...***




## Rereading Example: 3 page buffer

- Now, however, ***the smallest values are less than the largest (last) in the sorted run...***




# Repacking Example: 3 page buffer

- Once ***all buffer pages have a frozen value***, or input file is empty, start new run with the frozen values



# Repacking Example: 3 page buffer

- Once ***all buffer pages have a frozen value***, or input file is empty, start new run with the frozen values



# Repacking

- Note that, for buffer with  $B+1$  pages:
  - If input file is sorted → nothing is frozen → we get a **single run!**
  - If input file is reverse sorted (worst case) → everything is frozen → we get runs of length  **$B+1$**
- In general, with repacking we do **no worse** than without it!
- What if the file is already sorted?
- Engineer's approximation: runs will have  $\sim 2(B+1)$  length

$$\sim 2N \left( \left\lceil \log_B \frac{N}{2(B+1)} \right\rceil + 1 \right)$$

# Summary

- Basics of IO and buffer management.
  - See notebook for more fun! (Learn about *sequential flooding*)
- We introduced the IO cost model using **sorting**.
  - Saw how to do merges with few IOs,
  - Works better than main-memory sort algorithms.
- Described a few optimizations for sorting